

Ahmed Raza

AI Engineer | Full Stack AI Engineer | AI Software Developer

Karachi, Pakistan | ahmedrz.amjad@gmail.com | +92 330 7051591

[linkedin.com/in/ahmedraza101](https://www.linkedin.com/in/ahmedraza101) | github.com/ahmedraza-96 | ahmed-portfolio-jet.vercel.app

Summary

AI Engineer (BS in Artificial Intelligence, FAST-NUCES) building production LLM systems – RAG pipelines, tool-using agents, and evaluation pipelines – shipped end to end with FastAPI, Next.js, and AWS. Multiple products live.

Technical Skills

Generative AI: LLM APIs (OpenAI, Claude, Gemini, Groq), LangChain, LangGraph, RAG pipelines, AI agents, tool calling, LLM evals (LLM-as-judge, CI quality gates), Chroma, FAISS, LangSmith

Engineering: Python, TypeScript, SQL; FastAPI, Next.js, React, Node.js; PostgreSQL, MongoDB

Cloud & DevOps: AWS (EC2, S3, Lambda), Docker, Vercel, Terraform, GitHub Actions, Linux, Git

Experience

AI Engineer, HashOne Global – Karachi, Pakistan Jul 2025 – Jun 2026

- Built and shipped AI chatbots and tool-using agents end to end with LLM APIs and local LLMs.
- Designed RAG pipelines – chunking, embedding, retrieval tuning – to ground answers in source documents.
- Served LangGraph agents through FastAPI APIs; deployed with Docker and CI/CD on AWS.

AI Engineer Intern, Tech Hunt – Karachi, Pakistan Jun 2024 – Aug 2024

- Built RAG document pipelines in Python with LangChain, the OpenAI API, and FAISS.

Projects

NovaNet Eval Harness – AI support agent with production-grade LLM evals Live | Demo | Code

FastAPI, LangGraph, Groq Llama 3.3, Chroma, LLM-as-judge, GitHub Actions CI, Next.js, AWS EC2, Vercel

A customer-support RAG agent wrapped in a full evaluation pipeline: 100-case golden dataset, deterministic assertions plus LLM-as-judge grading, response caching, a CI quality gate that blocks prompt regressions before merge, and a Next.js metrics dashboard. Baseline evals caught an intermittent prompt-injection vulnerability.

PSX MCP Server – Published Model Context Protocol server for LLM agents PyPI | Demo | Code

Python, MCP (FastMCP), httpx, BeautifulSoup, pytest, GitHub Actions CI, PyPI trusted publishing

An open-source MCP server that gives Claude and other LLM agents live Pakistan Stock Exchange data – quotes, indices (KSE-100), company fundamentals, dividends, and announcements – as 10 tool-callable endpoints. Built test-first (78 tests), with CI across Python 3.10–3.13 on Linux and Windows, and published to PyPI for one-line install (`uvx psx-mcp-server`).

Convey AI – Conversational AI interview agent (Final Year Project) Live | Demo | Code

Next.js, Node.js, OpenAI API, REST APIs

An AI agent that conducts and evaluates mock interviews autonomously in Urdu and English, with an avatar-based interview mode. Deployed and live at conveyai.live.

Cognos – AI agent app (RAG + tools) Live | Demo | Code

FastAPI, LangChain / LangGraph, Chroma, PostgreSQL, Next.js, JWT, SSE, Terraform, AWS EC2

Lets users chat with a tool-using AI agent over their own documents, with built-in web search. Features JWT authentication, per-user Chroma vector stores, Postgres-backed chat memory, and token-by-token response streaming via SSE. Deployed on AWS EC2, provisioned with Terraform.

CareShare – Medicine donation platform (MERN) Live | Demo | Code

React, Node.js, Express, MongoDB

Connects medicine donors with people in need. Includes an admin panel that moderates donations and user accounts before listings go live. Deployed and accessible online.

IT Inventory Management System Live | Demo | Code

FastAPI, MySQL, React.js, REST APIs

Internal tool for tracking company IT assets and maintenance, with a React dashboard on a FastAPI and MySQL backend. Deployed and accessible online.

Education

Bachelor of Science in Artificial Intelligence (BSAI), FAST National University (NUCES), Karachi 2022 – 2026

Final Year Project: Convey AI – conversational AI interview agent (see Projects).