

Ahmed Raza

Full-Stack AI Engineer

Karachi, Pakistan | ahmedrz.amjad@gmail.com | +92 330 7051591

linkedin.com/in/ahmedraza101 | github.com/ahmedraza-96

Summary

Full-Stack AI Engineer with a BS in Artificial Intelligence from FAST-NUCES. Builds and ships complete web applications with Next.js, React, and FastAPI, plus production AI features – RAG pipelines, chatbots, and tool-using agents – using LLM APIs and local LLMs. Deploys with Docker, AWS, and Vercel; multiple products live.

Technical Skills

Frontend & Backend: Next.js, React.js; FastAPI, Flask, Node.js, Express; REST APIs, webhooks, JWT auth

Generative AI: LLMs, prompt engineering, LLM APIs (OpenAI, Anthropic Claude, Gemini), LangChain, LangGraph, RAG pipelines, AI agents, tool use / function calling, vector databases (Chroma, Qdrant, FAISS), LangSmith, Langfuse

Cloud, DevOps & Data: Docker, AWS (EC2, S3, Lambda), Vercel, Terraform, CI/CD (GitHub Actions), Linux, Git, AI coding agents & IDEs (Claude Code); PostgreSQL, Supabase, MongoDB, MySQL

Languages & ML: Python, JavaScript, TypeScript, SQL; pandas, NumPy, PyTorch, TensorFlow, scikit-learn, Keras

Experience

Full-Stack AI Engineer, HashOne Global – Karachi, Pakistan

Jul 2025 – Jun 2026

- Built AI applications and chatbots end to end using LLM APIs and local LLMs, applying prompt engineering to produce clean, structured output.
- Designed RAG pipelines over vector databases – document chunking, embedding, and retrieval tuning – to ground answers in source documents.
- Developed AI agents with LangChain and LangGraph that use tool calling to execute multi-step tasks, served through FastAPI REST APIs with webhook integrations.
- Deployed services with Docker and CI/CD pipelines on AWS, and monitored model quality with LangSmith.

AI Engineer Intern, Tech Hunt

Jun 2024 – Aug 2024

- Built RAG and document-processing pipelines in Python with LangChain and the OpenAI API.
- Chunked and embedded documents into a FAISS vector store and tuned retrieval for grounded answers.

Projects

Convey AI – Conversational AI interview agent (Final Year Project)

[Live link](#) | [Frontend code](#) | [Backend code](#)

Next.js, Node.js, OpenAI API, REST APIs

An AI agent that conducts and evaluates mock interviews autonomously in Urdu and English, with an avatar-based interview mode. Deployed and live at conveyai.live.

Cognos – Full-stack AI agent app (RAG + tools)

[Live link](#) | [Frontend code](#) | [Backend code](#)

FastAPI, LangChain / LangGraph, Chroma, PostgreSQL, Next.js, JWT, SSE, Terraform, AWS EC2

Lets users chat with a tool-using AI agent over their own documents, with built-in web search. Features JWT authentication, per-user Chroma vector stores, Postgres-backed chat memory, and token-by-token response streaming via SSE. Deployed on AWS EC2, provisioned with Terraform.

CareShare – Medicine donation platform (MERN)

[Live link](#) | [Frontend code](#) | [Backend code](#)

React, Node.js, Express, MongoDB

Connects medicine donors with people in need. Includes an admin panel that moderates donations and user accounts before listings go live. Deployed and accessible online.

IT Inventory Management System

[Live link](#) | [Frontend code](#) | [Backend code](#)

FastAPI, MySQL, React.js, REST APIs

Internal tool for tracking company IT assets and maintenance, with a React dashboard on a FastAPI and MySQL backend. Deployed and accessible online.

RAG Document Parsing Chatbot

Private client project

Python, LangChain, OpenAI, FAISS vector DB, FastAPI

A chatbot that answers questions over long documents – embeds the text into a FAISS vector store, retrieves the closest matches, and serves grounded answers through a FastAPI service.

Education

Bachelor of Science in Artificial Intelligence (BSAI), FAST National University (NUCES), Karachi

2022 – 2026

Relevant Coursework: Artificial Intelligence, Machine Learning, Deep Learning, Natural Language Processing, Computer Vision, Data Structures & Algorithms, Databases, Software Engineering

Final Year Project: Convey AI – conversational AI interview agent (see Projects).

Focus: Applied machine learning, LLMs, and building and shipping real AI products end to end.